

# Parallel Python for Data Processing

...

Embracing the Embarrassingly Parallel

# About Me

- Worked with Data since 2009, user of Python since 2011
- Live in Calgary (since 2014) with my wife and 2 year old daughter
- Developed and consulted on Research, Backend Services and Enterprise Data/AI use cases



**MSc in Bioinformatics  
- UCalgary**

Building Big Data Simulations on  
HPC Infrastructure



**Data/AI Team  
- ATB Financial**

Working alongside Data  
Scientists to Build/Deploy Models



**Cloud Solutions Architect -  
Oracle For Research**

Working alongside Researchers  
migrating workloads to the cloud



**Senior MLE  
- Coursera**

Working Alongside ML Scientists  
to Build/Deploy Models

**Parallelization can be  
simple, it's not all mutexes  
and race conditions.**

# Why Parallelize?

[1] ✓ 0.0s

□ ▾

```
print("Preprocessing Data ...")  
df = preprocess(data)  
  
print("Proprocessing Complete.")
```

[2] ↻ 9m 45.0s

... Preprocessing Data ...

PROBLEMS

OUTPUT

DEBUG CONSOLE

TERMINAL

PORTS

JUPYTER

POSTMAN CONSOLE

# Types of Parallel Problems

## Communication Intensive

Characterized by tasks that require **frequent data exchange or coordination** .

Use frameworks like MPI and hardware like HPC/Infiniband.

Examples:

- Simulations (weather forecasting, molecular dynamics)
- Large Language Model Training

## Embarrassingly Parallel

Characterized by many **independent** tasks

Use frameworks like Mapreduce and pools of instances

Examples:

- Pre-processing data to build a training set for ML modelling
  - cleaning, transforming, aggregating
- Running multiple independent simulations to test the effect of different parameters

# Ways we parallelize embarrassingly parallel Python processing

1. Shell Scripting
2. Python Multiprocessing (and Multithreading)
3. PySpark
4. Dask
5. Polars

# Shell Scripting



# Shell Scripting

Use shell commands to launch multiple Python scripts concurrently, each processing a portion of the data.

## Python

```
import pandas as pd
import sys

def calculate_mean_price(city_csv):
    df = pd.read_csv(city_csv)
    mean_price = df['price'].mean()
    print(f"{city_csv}: {mean_price}")

if __name__ == "__main__":
    if len(sys.argv) != 2:
        print("missing argument")
        sys.exit(1)
    city_csv = sys.argv[1]
    calculate_mean_price(city_csv)
```

## Bash

```
for city in Calgary Edmonton Vancouver Toronto; do
    python mean_housing_price.py $city.csv &
done
```

# Shell Scripting

Use shell commands to launch multiple Python scripts concurrently, each processing a portion of the data.

## Advantages

- Easy to implement
- Ready-made to run on a an HPC cluster with schedulers like Slurm or PBS

## Disadvantages

- Extra steps to aggregate and further process results

# Multiprocessing

# Multiprocessing vs Multithreading in Python

## Threads

Lightweight execution units with a single process.  
They share the same memory space.

Concurrent, but not parallel, execution (blame the GIL)

Good fit for I/O-bound tasks with lots of waiting  
(that said, you should look at asyncio for performance)

## Processes

Independent execution environments with their own memory space and Python interpreter

True parallel execution (on multi-core processors)

Good fit for CPU-bound tasks

# Multiprocessing vs Multithreading in Python

```
import threading
import time

def task(name):
    print(f"Thread {name}: starting")
    time.sleep(2)
    print(f"Thread {name}: finishing")

if __name__ == "__main__":
    threads = []
    for i in range(3):
        thread = threading.Thread(target=task,
                                  args=(i,))
        threads.append(thread)
        thread.start()

    for thread in threads:
        thread.join()
```

```
from multiprocessing import Pool
import time

def task(name):
    print(f"Process {name}: starting")
    time.sleep(2)
    print(f"Process {name}: finishing")

if __name__ == "__main__":
    with Pool(processes=3) as pool:
        pool.map(task, range(3))
```

# Why Multiprocessing over Shell Scripting?

## Advantages

- Integrates seamlessly with the rest of your Python code -> can immediately process results within same python session
- Better error handling and debugging

## Disadvantages

- Development overhead for small tasks

# Threading and No-GIL

There are ongoing efforts to remove the GIL from CPython. This could significantly improve parallel performance for CPU-bound tasks using threads.

## What's New In Python 3.13

**Editors:** Adam Turner and Thomas Wouters

This article explains the new features in Python 3.13, compared to 3.12. Python 3.13 was released on October 7, 2024. For full details, see the [changelog](#).

**See also:** [PEP 719](#) – Python 3.13 Release Schedule

## Summary – Release Highlights

Python 3.13 is the latest stable release of the Python programming language, with a mix of changes to the language, the implementation and the standard library. The biggest changes include a new [interactive interpreter](#), experimental support for running in a [free-threaded mode](#) ([PEP 703](#)), and a [Just-In-Time compiler](#) ([PEP 744](#)).

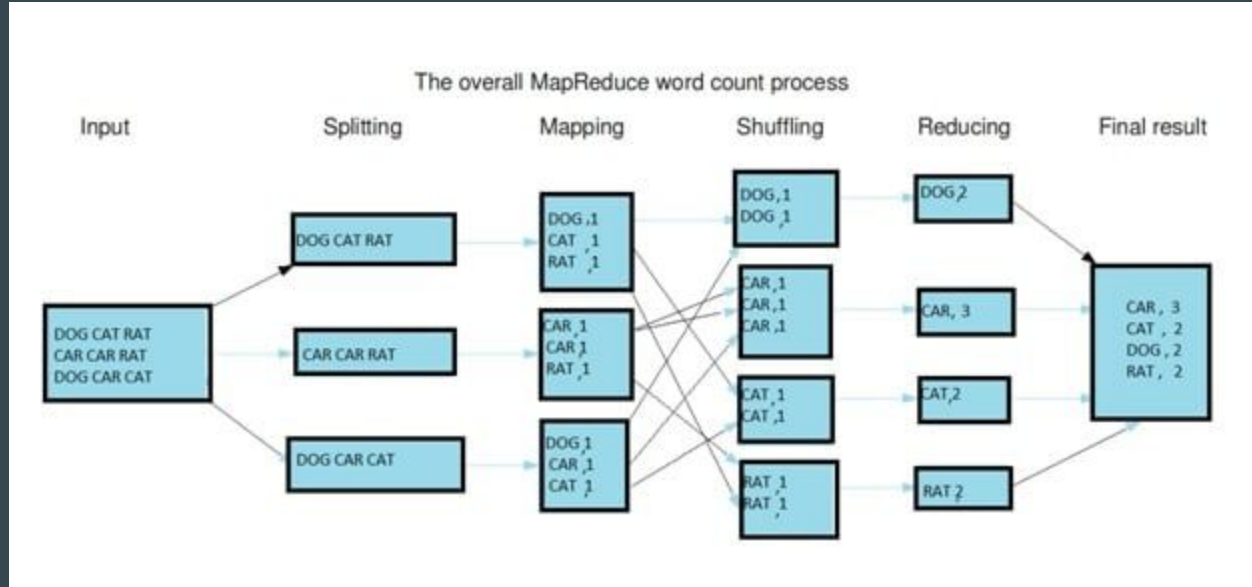
# A note on Jupyter and Multiprocessing

- The `if __name__ == '__main__':` Guard ensures code creating new processes is only run when a script is executed (not imported as a module).
  - Jupyter notebooks execute code in a way that confused the multiprocessing libraries, leading to issues if this guard is missing.
- Jupyter can also introduce issues with 'pickling' and global variables.
- When using jupyter, try `ipyparallel`



# PySpark

# MapReduce



# Spark

- Developed as a faster and more versatile alternative to Hadoop's MapReduce (but often still runs with Hadoop)
- Supports Batch Processing, Stream Processing, ML and Graph Processing
- Distributes data and computations across a cluster of machines
  - Process really big datasets

# PySpark

- PySpark is the Python API for Spark
- Write spark code with Python syntax (using dataframes similar to pandas, but distributed)

# Demo 1



# Our Sample Problem

## NYC Yellow Taxi Trip Data

Practice your ML skills on this Time-Series Dataset!



[Data Card](#) [Code \(14\)](#) [Discussion \(1\)](#) [Suggestions \(0\)](#)

### About Dataset

#### Context

New York City (NYC) Taxi & Limousine Commission (TLC) keeps data from all its cabs, and it is freely available to download from its official website. You can access it [here](#). Now, the TLC primarily keeps and manages data for 4 different types of vehicles:

- **Yellow Taxi: Yellow Medallion Taxicabs:** These are the famous NYC yellow taxis that provide transportation exclusively through street hails. The number of taxicabs is limited by a finite number of medallions issued by the TLC. You access this mode of transportation by standing in the street and hailing an available taxi with your hand. The pickups are not pre-arranged.
- **Green Taxi: Street Hail Livery:** The SHL program will allow livery vehicle owners to license and outfit their vehicles with green borough taxi branding, meters, credit card machines, and ultimately the right to accept street hails in addition to pre-arranged rides.
- **For-Hire Vehicles (FHVs):** FHV transportation is accessed by a pre-arrangement with a dispatcher or limo company. These FHVs are not permitted to pick up passengers via street hails, as those rides are not considered pre-arranged.

#### Usability ⓘ

10.00

#### License

[U.S. Government Works](#)

#### Expected update frequency

Never

#### Tags

Tabular

Travel

Time Series Analysis

Regression

Clustering

Dask

# Dask

An easy to use, flexible library for parallel computing in Python.

Scales from single machines to clusters.

Provides familiar APIs, mirroring NumPy and Pandas.

Breaks large computations into smaller tasks and builds a task graph.

Dynamically schedules and executes tasks based on dependencies and available resources.



Optimizes data locality and minimizes communication overhead.

# Dask Collections

Dataframes - Parallelize the pandas library

Arrays - Parallelize the NumPy library

- Larger-than-memory execution for single machines
- Parallel execution, Distributed when available





# Why Dask?

## Advantages

- Easy to use or add to your numpy/pandas workflow
- Flexible to various parallel computing needs
- Scales from 1 machine to many

## Disadvantages

- Scheduling/communication overhead, noticeable on smaller tasks

# Dask vs Spark

- Generally Dask is smaller and lighter weight than Spark.
  - Fewer features and, instead, is used in conjunction with other libraries
- Many enterprise data lakes are built on top of spark (e.g. Databricks, AWS Elastic MapReduce, GCP Dataproc, Cloudera, etc.) making it a default choice.
- [https://www.reddit.com/r/Python/comments/198f5vc/one\\_billion\\_row\\_challenge\\_dask\\_vs\\_spark/](https://www.reddit.com/r/Python/comments/198f5vc/one_billion_row_challenge_dask_vs_spark/)

# Polars

# Polars vs Dask

When your data fits in memory -> Try Polars first

Polars also *may* be better for text processing

# Demo 2



# ML Package Parallelism

TensorFlow - CPU/GPU/TPU Parallelism (but not always easy)

Scikit-learn - Some operations support CPU parallelism

XGBoost - Supports CPU/Multi-machine parallelism

LightGBM - CPU/GPU Parallelism

PyTorch - CPU/GPU/Multi-machine Parallelism

# How to choose

Have you verified you need parallelism

No -> KISS

Yes: Does your package of choice support it natively?

Yes -> Use that package

No: Do you have access to a public cloud or a cluster?

No -> Dask/Polars

Yes:

You're in academia with access to ARC -> Shell Scripting

You're in industry with access to Spark/Databricks -> PySpark

Other -> Dask

# Thanks!

<https://www.linkedin.com/in/nathanbryans>

[https://github.com/nbryans/parallel\\_python\\_examples](https://github.com/nbryans/parallel_python_examples)

```
Advent of Code  [About]  [Events]  [Shop]  [Log In]
0xffff&2024      [Calendar]  [AoC++]  [Sponsors]  [Leaderboard]  [Stats]
```

The first puzzles will unlock on December 1st at midnight EST (UTC-5). See you then!

In the meantime, you can still access past [Events].